

Peiyang Xu

Princeton University, NJ 08540 United States

(+1)640-230-4824 | px4668@princeton.edu

Website: <https://xupy2003.github.io/>

Education

Electrical and Computer Engineering, Princeton University

Princeton, NJ

Ph.D. student in Electrical and Computer Engineering

Sep 2025 – Now

Advisor: Prateek Mittal & Pramod Viswanath

Department of Computer Science and Technology, Tsinghua University

Beijing, China

Bachelor of Engineering in Computer Science and Technology

Sep 2021 – Jun 2025

- Courses: Software Engineering (A+), Introduction to Complex Analysis (A+), Student Research Training (A+), Modern Cryptography (A+), Computer Graphics (A), Artificial Neural Network (A)
- Research interest: AI Safety, Reinforcement Learning, Multimodal Models, Agentic systems

Research Experience

Project: Context-Aware RL for Agentic and Multimodal LLMs

Princeton, NJ

Collaborator: Xingyu Fu Advisor: Prateek Mittal, Pramod Viswanath

Dec 2025 – May 2026

- Identified “context unawareness” as a shared failure mode of agentic and multimodal LLMs, where models fail to ground predictions in sparse yet decisive evidence.
- Constructed contrastive context pairs across two domains: 1k agentic trajectory pairs via a cascade of condition filters, and 7k multimodal image pairs built through generative editing and similarity-based retrieval.
- Proposed ContextRL, a modality-agnostic post-training method that augments GRPO with an auxiliary context-selection objective, rewarding the policy through a clipped logit-level contrastive loss for choosing the context that supports a given query–answer pair over a minimally perturbed confounder.
- Achieved consistent gains over standard GRPO: +2.2% across 5 long-horizon agentic and long-context benchmarks and +1.8% across 12 diverse VQA benchmarks.
- Demonstrated through controlled data-augmentation baselines that reuse the same contrastive data but collapse or add nothing, isolating the gains to the context-selection objective rather than the data alone.

Project: Verifiable Process Supervision for Language Models

Princeton, NJ

Advisor: Pramod Viswanath, Sewoong Oh

Sep 2025 – Apr 2026

- Proposed Verifiable Process Supervision (VPS), a post-training framework for verifiable domains that jointly optimizes prediction accuracy and reasoning quality.
- Induced a structured reasoning prior via SFT on synthetically constructed traces, enabling reliable syntactic extraction and verification of intermediate claims without relying on expensive or noisy LLM judges.
- Introduced adaptive reward weighting that allocates greater weight to reasoning subtasks with the largest remaining errors, inducing an implicit curriculum over reasoning subskills.
- Showed on chess that accuracy-only RL degrades reasoning quality, whereas VPS preserves accuracy while reducing win-rate error by up to 30% and restoring consistency to near saturation.

Project: Mitigating Object Hallucination in Large Vision-Language Models through Adversarial Contrastive Finetuning

Beijing, China

Advisor: Xiaolin Hu, Associate Professor, Tsinghua University

Feb 2025 – Aug 2025

- Explored the internal mechanism of object hallucination from the perspective of image embedding based on cosine similarity and Smooth Grad-CAM
- Proposed an adversarial contrast fine-tuning (ACFT) method to enhance the alignment between visual and textual embedding, thus mitigating object hallucination
- Provided a pipeline to automatically generating paired positive and negative examples using an adversarial hallucination attribute flipping (AHAF) method, followed by contrastive fine-tuning of the LVLM
- Achieved state-of-the-art performance on multiple benchmarks, e.g. outperforming existing approaches like VCD, OPERA and VTI, etc. on benchmarks like POPE and MME

Project: SafeVision: Efficient Image Guardrail with Robust Policy Adherence and Explainability

Chicago, IL

Advisor: Bo Li, Associate Professor, Dept of CS, University of Chicago

July 2024 – Nov 2024

- Proposed SafeVision, a novel image guardrail system that integrates human-like understanding and reasoning, with robust policy adherence, explainability, and lightning-fast inference speeds
- Provided VisionHARM-500K, a high-quality unsafe image benchmark comprising over 500k images to cover a wide array of risky categories, significantly broadening the scope and depth of unsafe image benchmarks
- Implemented an effective data collection and generation framework, a policy-following training pipeline, a customized loss function, and an efficient diverse QA generation strategy to enhance the training effectiveness

- Achieved state-of-the-art performance in both efficiency and accuracy, with an accuracy of 91.36% on VisionHARM-500K (17.36% higher than GPT-4o) and an inference time of 0.313 seconds per image

Project: MMDT: Decoding the Trustworthiness and Safety of Multimodal Foundation Models Beijing, China
 Advisor: Bo Li, Associate Professor, Dept of CS, University of Chicago Feb 2024 – Jun 2024

- Aimed to provide the first comprehensive safety and trustworthiness evaluation dataset and platform for multimodal foundation models (MMFMs), to help develop safer and more reliable MMFMs and systems
- Designed various evaluation scenarios/tasks and red teaming (attack) algorithms to generate challenging data and form a high-quality benchmark, to assess models from multiple perspectives, including safety, hallucination, fairness/bias, privacy, adversarial robustness, and out-of-distribution (OOD)
- Evaluated a range of state-of-the-art multimodal models, such as GPT-4o, Gemini, and Qwen, using our MMDT platform, and revealed several vulnerabilities and areas for improvement across all perspectives for these models

Project: Natural Language Induced Adversarial Images Beijing, China
 Advisor: Xiaolin Hu, Associate Professor, Tsinghua University July 2023 – April 2024

- Aimed to use prompt optimization to perform black-box attacks on commercial text-to-image models such as Midjourney and DALL-E-3, by generating natural adversarial examples in order to let classifiers make mistakes
- Utilized and improved non-differentiable optimization methods such as Genetic Algorithm and Ant Colony Algorithm to optimize the input prompt in order to achieve the attacking purpose
- Performed Genetic Algorithm in the discrete prompt space and assessed the fitness of individuals using a weighted result of attack success rate and image naturalness, with the naturalness measured by CLIP
- First to explore using prompt-based optimization strategies for generating adversarial images, with an attack success rate of 92.1% and superior in terms of naturalness, transitivity, and interpretability
- Analyzed adversarial images from a natural language perspective and demonstrated that the adversarial semantic information extracted from these images significantly affects the accuracy of traditional image classifiers

Project: CurBench: Curriculum Learning Benchmark Beijing, China
 Advisor: Xin Wang, Assistant Researcher, Tsinghua University Oct 2022 – Jan 2023

- Aimed to use curriculum learning to improve various classification algorithms and to establish the first benchmark for a systematic evaluation on curriculum learning
- Created a benchmark consisting of 16 datasets spanning 3 research domains (computer vision, natural language processing, and graph machine learning) under 3 settings (standard, noised, and imbalanced)
- Provided a unified pipeline that plugs automatic curricula into general machine learning process, supporting 14 core curriculum learning methods (personally responsible for the Meta Weight Net)

Publications & Manuscripts

- Peiyang Xu**, Bangzheng Li, Sijia Liu, Karthik R. Narasimhan, Pramod Viswanath, Prateek Mittal, Xingyu Fu: [Context-Aware RL for Agentic and Multimodal LLMs](#) (submitted to NeurIPS 2026, *under review*)
- Kyuyoung Kim, Kevin Wang, Yunfei Xie, **Peiyang Xu**, Peiyao Sheng, Chen Wei, Zhangyang Wang, Jinwoo Shin, Pramod Viswanath, Sewoong Oh: [Correct Answers from Sound Reasoning: Verifiable Process Supervision for Language Models](#) (submitted to NeurIPS 2026, *under review*)
- Peiyang Xu**, Xiaopei Zhu, Jun Zhu, Xiaolin Hu: Mitigating Object Hallucination in Large Vision-Language Models through Adversarial Contrastive Finetuning (submitted to ACM MM 2026, *under review*)
- Peiyang Xu**, Minzhou Pan, Zhaorun Chen, Shuang Yang, Chaowei Xiao, Bo Li, [SafeVision: Efficient Image Guardrail with Robust Policy Adherence and Explainability](#)
- Chejian Xu, Jiawei Zhang, Zhaorun Chen, Chulin Xie, Mintong Kang, Zhuowen Yuan, Zidi Xiong, Chenhui Zhang, Lingzhi Yuan, Yi Zeng, **Peiyang Xu**, Chengquan Guo, Andy Zhou, Jeffrey Zhiwei Tan, Zhen Xiang, Zinan Lin, Dan Hendrycks, Dawn Song, Bo Li, [MMDT: Decoding the Trustworthiness and Safety of Multimodal Foundation Models](#) (ICLR 2025)
- Xiaopei Zhu*, **Peiyang Xu***, Guanning Zeng, Yinpeng Dong, Xiaolin Hu, [Natural Language Induced Adversarial Images](#) (ACM MM 2024).
- Yuwei Zhou, Zirui Pan, Xin Wang, Hong Chen, Haoyang Li, Yanwen Huang, Zhixiao Xiong, Fangzhou Xiong, **Peiyang Xu**, Shengnan Liu and Wenwu Zhu, [CurBench: Curriculum Learning Benchmark](#), accepted by 42th International Conference on Machine Learning (ICML 2024)

Honors & Awards

- Gordon Y.S. Wu Fellowship in Engineering, Princeton University(2025)
- Outstanding Graduate, Department of Computer Science and Technology, Tsinghua University(2025)
- Tsinghua University Research Excellence Scholarship (2024)
- Tsinghua University Academic Excellence Scholarship (2024)

Skills & Others

Language: Chinese (Native); English (Fluent)

Exams: TOFEL: 108; GRE: 158+170+4.0

Computer Science: Proficient in Python, C++, and various other programming languages; Deep understanding of multimodal language models